

adult-income-dataset-intel-5

July 8, 2024

0.0.1 Loading Necessary Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.linear_model import LogisticRegression

# Warning Filter
import warnings
warnings.filterwarnings('ignore')
```

0.0.2 Loading Dataset

```
[2]: df = pd.read_csv('/content/adult.csv')
df.head(10)
```

```
[2]:
```

	age	workclass	fnlwgt	education	educational-num	\
0	25	Private	226802	11th	7	
1	38	Private	89814	HS-grad	9	
2	28	Local-gov	336951	Assoc-acdm	12	
3	44	Private	160323	Some-college	10	
4	18	?	103497	Some-college	10	
5	34	Private	198693	10th	6	
6	29	?	227026	HS-grad	9	
7	63	Self-emp-not-inc	104626	Prof-school	15	
8	24	Private	369667	Some-college	10	
9	55	Private	104996	7th-8th	4	

	marital-status	occupation	relationship	race	gender	\
0	Never-married	Machine-op-inspct	Own-child	Black	Male	
1	Married-civ-spouse	Farming-fishing	Husband	White	Male	
2	Married-civ-spouse	Protective-serv	Husband	White	Male	
3	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	
4	Never-married	?	Own-child	White	Female	

5	Never-married	Other-service	Not-in-family	White	Male
6	Never-married	?	Unmarried	Black	Male
7	Married-civ-spouse	Prof-specialty	Husband	White	Male
8	Never-married	Other-service	Unmarried	White	Female
9	Married-civ-spouse	Craft-repair	Husband	White	Male

	capital-gain	capital-loss	hours-per-week	native-country	income
0	0	0	40	United-States	<=50K
1	0	0	50	United-States	<=50K
2	0	0	40	United-States	>50K
3	7688	0	40	United-States	>50K
4	0	0	30	United-States	<=50K
5	0	0	30	United-States	<=50K
6	0	0	40	United-States	<=50K
7	3103	0	32	United-States	>50K
8	0	0	40	United-States	<=50K
9	0	0	10	United-States	<=50K

0.0.3 Learning More about Dataset

```
[3]: df.columns
```

```
[3]: Index(['age', 'workclass', 'fnlwgt', 'education', 'educational-num',
        'marital-status', 'occupation', 'relationship', 'race', 'gender',
        'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
        'income'],
        dtype='object')
```

```
[4]: df.shape
```

```
[4]: (48842, 15)
```

```
[5]: df.dtypes
```

```
[5]: age                int64
workclass             object
fnlwgt               int64
education            object
educational-num      int64
marital-status       object
occupation           object
relationship         object
race                object
gender              object
capital-gain         int64
capital-loss         int64
hours-per-week       int64
```

```
native-country    object
income            object
dtype: object
```

```
[6]: df.isnull().sum()
```

```
[6]: age                0
workclass             0
fnlwgt               0
education             0
educational-num       0
marital-status        0
occupation            0
relationship          0
race                 0
gender               0
capital-gain          0
capital-loss          0
hours-per-week        0
native-country        0
income               0
dtype: int64
```

```
[7]: df.nunique()
```

```
[7]: age                74
workclass              9
fnlwgt               28523
education             16
educational-num       16
marital-status         7
occupation            15
relationship           6
race                   5
gender                 2
capital-gain          123
capital-loss           99
hours-per-week         96
native-country         42
income                 2
dtype: int64
```

```
[8]: df.describe().T
```

```
[8]:
```

	count	mean	std	min	25%	\
age	48842.0	38.643585	13.710510	17.0	28.0	
fnlwgt	48842.0	189664.134597	105604.025423	12285.0	117550.5	

educational-num	48842.0	10.078089	2.570973	1.0	9.0
capital-gain	48842.0	1079.067626	7452.019058	0.0	0.0
capital-loss	48842.0	87.502314	403.004552	0.0	0.0
hours-per-week	48842.0	40.422382	12.391444	1.0	40.0

	50%	75%	max
age	37.0	48.0	90.0
fnlwgt	178144.5	237642.0	1490400.0
educational-num	10.0	12.0	16.0
capital-gain	0.0	0.0	99999.0
capital-loss	0.0	0.0	4356.0
hours-per-week	40.0	45.0	99.0

0.0.4 Value Count Function

```
[9]: df['workclass'].value_counts() #fillna
```

```
[9]: workclass
Private          33906
Self-emp-not-inc  3862
Local-gov        3136
?                2799
State-gov        1981
Self-emp-inc     1695
Federal-gov      1432
Without-pay       21
Never-worked      10
Name: count, dtype: int64
```

```
[10]: df['occupation'].value_counts() #fillna
```

```
[10]: occupation
Prof-specialty    6172
Craft-repair      6112
Exec-managerial   6086
Adm-clerical      5611
Sales             5504
Other-service     4923
Machine-op-inspct 3022
?                2809
Transport-moving  2355
Handlers-cleaners 2072
Farming-fishing   1490
Tech-support      1446
Protective-serv   983
Priv-house-serv   242
Armed-Forces      15
```

Name: count, dtype: int64

```
[11]: df['native-country'].value_counts() #fillna
```

```
[11]: native-country
United-States      43832
Mexico             951
?                  857
Philippines        295
Germany            206
Puerto-Rico       184
Canada             182
El-Salvador        155
India              151
Cuba               138
England            127
China              122
South              115
Jamaica            106
Italy              105
Dominican-Republic 103
Japan              92
Guatemala          88
Poland             87
Vietnam            86
Columbia           85
Haiti              75
Portugal           67
Taiwan             65
Iran               59
Greece             49
Nicaragua          49
Peru               46
Ecuador            45
France             38
Ireland            37
Hong               30
Thailand           30
Cambodia           28
Trinidad&Tobago    27
Laos               23
Yugoslavia         23
Outlying-US(Guam-USVI-etc) 23
Scotland           21
Honduras           20
Hungary            19
Holand-Netherlands 1
```

Name: count, dtype: int64

```
[12]: df['income'].value_counts() #0,1
```

```
[12]: income
<=50K    37155
>50K     11687
Name: count, dtype: int64
```

```
[13]: df['marital-status'].value_counts()
```

```
[13]: marital-status
Married-civ-spouse    22379
Never-married         16117
Divorced               6633
Separated             1530
Widowed               1518
Married-spouse-absent   628
Married-AF-spouse       37
Name: count, dtype: int64
```

```
[14]: df['gender'].value_counts()
```

```
[14]: gender
Male    32650
Female  16192
Name: count, dtype: int64
```

```
[15]: df['race'].value_counts()
```

```
[15]: race
White    41762
Black    4685
Asian-Pac-Islander  1519
Amer-Indian-Eskimo   470
Other             406
Name: count, dtype: int64
```

0.0.5 Removing rows with ? Values

```
[16]: # Replacing ? with NaN
df['workclass'] = df['workclass'].replace('?', np.nan)
df['occupation'] = df['occupation'].replace('?', np.nan)
df['native-country'] = df['native-country'].replace('?', np.nan)

# Drop rows with any NaN values
df = df.dropna()
```

```
# Check the shape of the DataFrame
df.shape

df2=df.copy()
```

0.0.6 Feature Engineering

```
[17]: # For Education
df.education = df.education.
    ↪replace(['Preschool','1st-4th','5th-6th','7th-8th','9th','10th','11th','12th'],'School')
df.education = df.education.replace('HS-grad','High School')
df.education = df.education.
    ↪replace(['Assoc-voc','Assoc-acdm','Prof-school','Some-college'],'Higher-Education')
df.education = df.education.replace('Bachelors','Under-Grad')
df.education = df.education.replace('Masters','Graduation')
df.education = df.education.replace('Doctorate','Doc')
```

```
[18]: # For Marital status
df['marital-status'] = df['marital-status'].
    ↪replace(['Married-civ-spouse','Married-AF-spouse'],'Married')
df['marital-status'] = df['marital-status'].
    ↪replace(['Never-married'],'Unmarried')
df['marital-status'] = df['marital-status'].
    ↪replace(['Divorced','Separated','Widowed','Married-spouse-absent'],'Single')
```

```
[19]: # For Income
df['income'] = df['income'].replace({'<=50K': 0, '>50K': 1})
```

```
[20]: df
```

```
[20]:
```

	age	workclass	fnlwgt	education	educational-num	\
0	25	Private	226802	School	7	
1	38	Private	89814	High School	9	
2	28	Local-gov	336951	Higher-Education	12	
3	44	Private	160323	Higher-Education	10	
5	34	Private	198693	School	6	
...	
48837	27	Private	257302	Higher-Education	12	
48838	40	Private	154374	High School	9	
48839	58	Private	151910	High School	9	
48840	22	Private	201490	High School	9	
48841	52	Self-emp-inc	287927	High School	9	
	marital-status	occupation	relationship	race	gender	\
0	Unmarried	Machine-op-inspct	Own-child	Black	Male	

1	Married	Farming-fishing	Husband	White	Male
2	Married	Protective-serv	Husband	White	Male
3	Married	Machine-op-inspct	Husband	Black	Male
5	Unmarried	Other-service	Not-in-family	White	Male
...
48837	Married	Tech-support	Wife	White	Female
48838	Married	Machine-op-inspct	Husband	White	Male
48839	Single	Adm-clerical	Unmarried	White	Female
48840	Unmarried	Adm-clerical	Own-child	White	Male
48841	Married	Exec-managerial	Wife	White	Female

	capital-gain	capital-loss	hours-per-week	native-country	income
0	0	0	40	United-States	0
1	0	0	50	United-States	0
2	0	0	40	United-States	1
3	7688	0	40	United-States	1
5	0	0	30	United-States	0
...
48837	0	0	38	United-States	0
48838	0	0	40	United-States	1
48839	0	0	40	United-States	0
48840	0	0	20	United-States	0
48841	15024	0	40	United-States	1

[45222 rows x 15 columns]

[21]: df2

[21]:

	age	workclass	fnlwgt	education	educational-num	\
0	25	Private	226802	11th	7	
1	38	Private	89814	HS-grad	9	
2	28	Local-gov	336951	Assoc-acdm	12	
3	44	Private	160323	Some-college	10	
5	34	Private	198693	10th	6	
...	
48837	27	Private	257302	Assoc-acdm	12	
48838	40	Private	154374	HS-grad	9	
48839	58	Private	151910	HS-grad	9	
48840	22	Private	201490	HS-grad	9	
48841	52	Self-emp-inc	287927	HS-grad	9	

	marital-status	occupation	relationship	race	gender	\
0	Never-married	Machine-op-inspct	Own-child	Black	Male	
1	Married-civ-spouse	Farming-fishing	Husband	White	Male	
2	Married-civ-spouse	Protective-serv	Husband	White	Male	
3	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	
5	Never-married	Other-service	Not-in-family	White	Male	

...
48837	Married-civ-spouse		Tech-support		Wife	White Female
48838	Married-civ-spouse		Machine-op-inspct		Husband	White Male
48839		Widowed	Adm-clerical		Unmarried	White Female
48840		Never-married	Adm-clerical		Own-child	White Male
48841	Married-civ-spouse		Exec-managerial		Wife	White Female

	capital-gain	capital-loss	hours-per-week	native-country	income
0	0	0	40	United-States	<=50K
1	0	0	50	United-States	<=50K
2	0	0	40	United-States	>50K
3	7688	0	40	United-States	>50K
5	0	0	30	United-States	<=50K

...
48837	0	0	38	United-States	<=50K	
48838	0	0	40	United-States	>50K	
48839	0	0	40	United-States	<=50K	
48840	0	0	20	United-States	<=50K	
48841	15024	0	40	United-States	>50K	

[45222 rows x 15 columns]


0.1 2. df2= Pre Column Reductions (Used in Plots)

0.1.1 Some Graphs and Plots

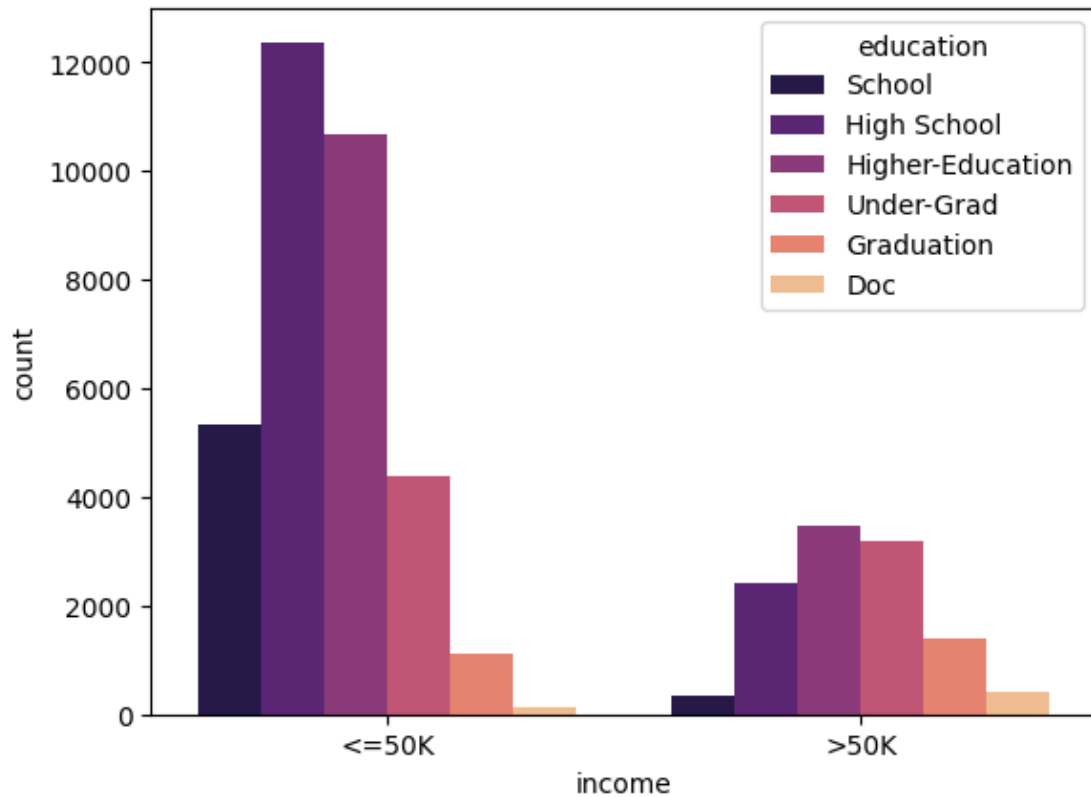
```
[22]: # Select only numeric columns
numeric_df = df.select_dtypes(include=['number'])

# Calculate the correlation matrix for numeric columns
correlation_matrix = numeric_df.corr()


# Display the heatmap
sns.heatmap(correlation_matrix, annot=True);
```



```
[23]: # @title Income with Education
sns.countplot(x=df2['income'],palette='magma',hue='education',data=df);
```




```
[24]: # @title Income Distribution by Workclass
plt.figure(figsize=(14,5))
sns.histplot(data=df2, x="workclass", hue="income",multiple="stack")
plt.title("Income Distribution by Workclass")
plt.xlabel("Workclass")
plt.ylabel("Count")
plt.show()
print(df2['workclass'].unique())
```




```
['Private' 'Local-gov' 'Self-emp-not-inc' 'Federal-gov' 'State-gov'
 'Self-emp-inc' 'Without-pay']
```

```
[25]: # @title Income Distribution by Workclass
plt.figure(figsize = (15,5))
sns.countplot(data = df2, x = 'income',palette='magma', hue = 'workclass');
```




```
[26]: # @title Occupation Distribution by Income
plt.figure(figsize=(28,5))
sns.countplot(data=df2, x="occupation", hue="income")
plt.title("Occupation Distribution by Income")
plt.xlabel("Occupation")
plt.ylabel("Count")
plt.show()
print(df2['occupation'].unique())
```




```
['Machine-op-inspct' 'Farming-fishing' 'Protective-serv' 'Other-service'
 'Prof-specialty' 'Craft-repair' 'Adm-clerical' 'Exec-managerial'
 'Tech-support' 'Sales' 'Priv-house-serv' 'Transport-moving'
 'Handlers-cleaners' 'Armed-Forces']
```

```
[27]: # @title Education Graph 2
df.groupby('education').size().plot(kind='barh', color=sns.palettes.
      ↪mpl_palette('Dark2'))
plt.gca().spines[['top', 'right',]].set_visible(False)
```




```
[28]: # @title Income Distribution by Marital Status
plt.figure(figsize=(18,5))
sns.countplot(data=df2, x="marital-status", hue="income")
plt.title("Income Distribution by Marital Status")
plt.xlabel("Marital Status")
plt.ylabel("Count")
plt.show()
print(df2['marital-status'].unique())
```



```
['Never-married' 'Married-civ-spouse' 'Widowed' 'Separated' 'Divorced'  
 'Married-spouse-absent' 'Married-AF-spouse']
```


```
[29]: # @title Income Graph  
#df2 for income plots  
df2['income'].value_counts().plot(kind='bar', color=sns.color_palette('Dark2'))  
plt.title('Distribution of Income')  
plt.xlabel('Income')  
plt.ylabel('Count')  
plt.xticks(rotation=0)  
plt.show()
```



```
[30]: # @title Age Distribution by Income  
  
# Plot the histogram with a custom hue using Seaborn  
plt.figure(figsize=(10, 3))  
sns.histplot(data=df2, x="age", hue="income", kde=True, multiple="layer")  
plt.title("Age Distribution by Income")
```

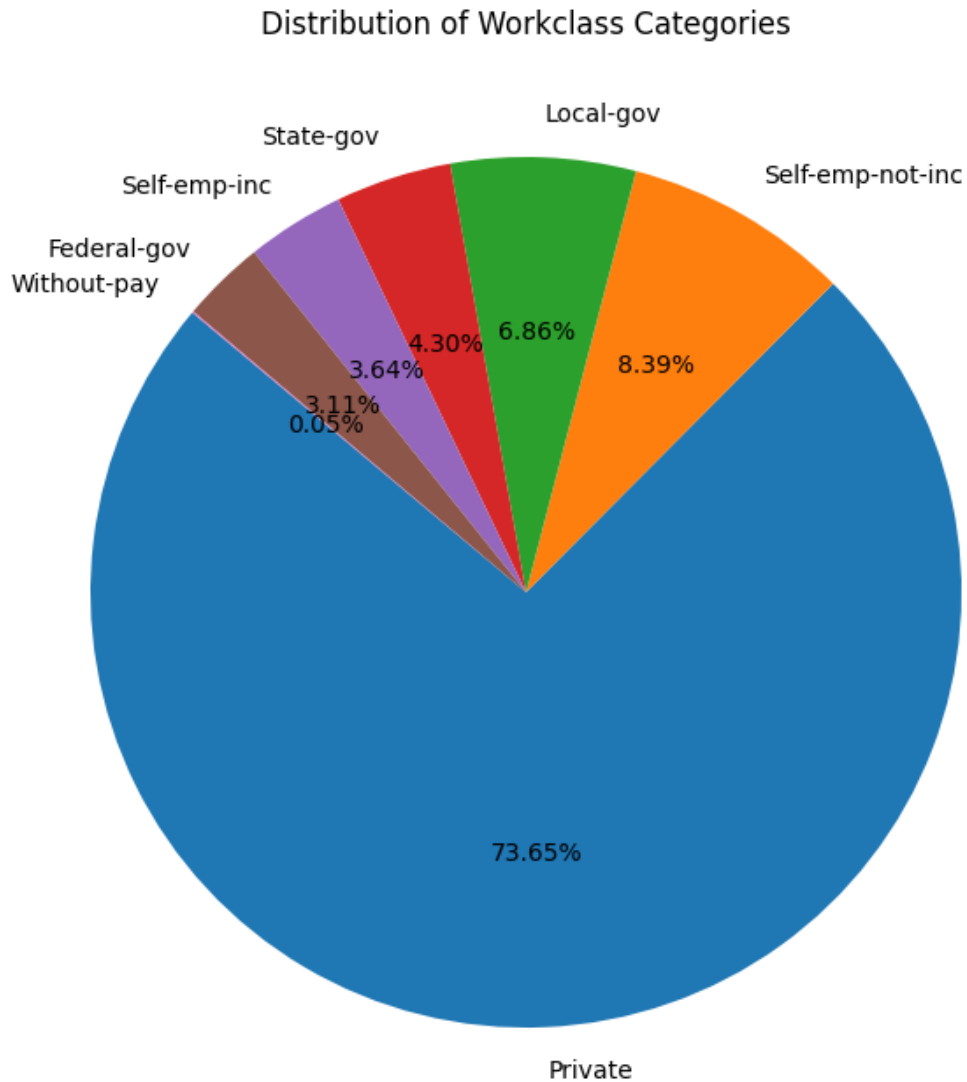
```
plt.xlabel("Age")
plt.ylabel("Count")
plt.show()

# Get the sorted unique ages
unique_age = df2['age'].unique()
sorted_age = np.sort(unique_age)
print(sorted_age)
```




```
[17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64
 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88
 89 90]
```

```
[31]: # @title Distribution of Workclass Categories
workclass_counts = df2['workclass'].value_counts()
labels = workclass_counts.index
sizes = workclass_counts.values
plt.figure(figsize=(8,14))
plt.pie(sizes,labels=labels, autopct='%1.2f%%', startangle=140)
plt.title("Distribution of Workclass Categories")
plt.show()
print(df2['workclass'].unique())
```




```
[ 'Private' 'Local-gov' 'Self-emp-not-inc' 'Federal-gov' 'State-gov'
  'Self-emp-inc' 'Without-pay']
```

```
[32]: # @title Income Distribution by Gender
plt.figure(figsize=(4,3))
sns.countplot(data=df2,x="gender",hue="income")
plt.title("Income Distribution by Gender")
plt.xlabel("Gender")
plt.ylabel("Count")
plt.legend()
plt.show()
print(df2['gender'].unique())
```

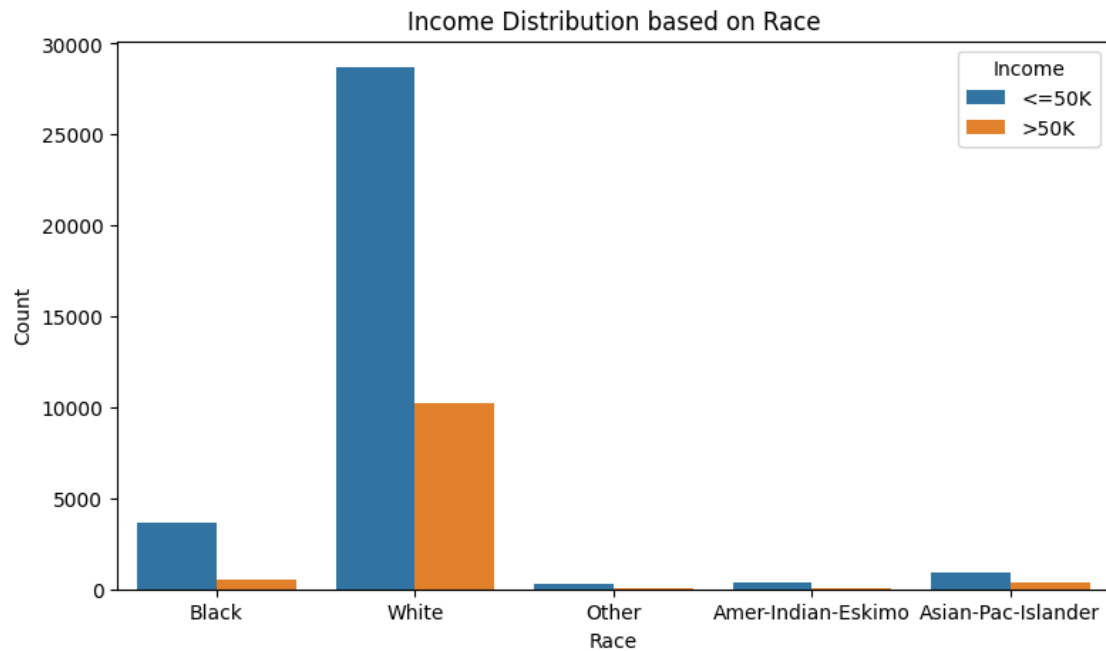



['Male' 'Female']

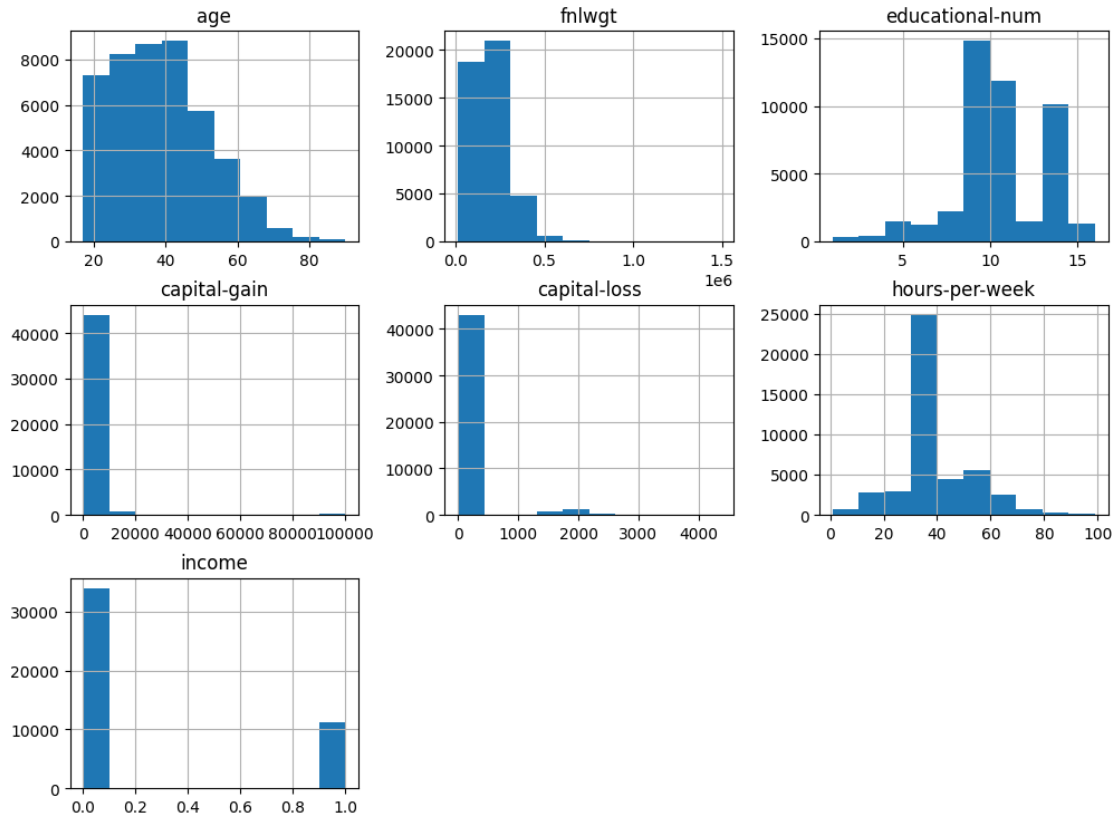
```
[33]: # @title Income Distribution by Education
education_order = df2.sort_values('educational-num')['education'].unique()
plt.figure(figsize=(7,3))
sns.countplot(data=df2,x="education",hue="income",order=education_order)
plt.title("Income Distribution by Education Level")
plt.xlabel("Education Level")
plt.ylabel("Count")
plt.xticks(rotation=60)
plt.legend(title="Income")
plt.show()
```




```
[36]: # @title Income Distribution based on Race
plt.figure(figsize=(9, 5))
sns.countplot(data=df2, x="race", hue="income")
plt.title("Income Distribution based on Race")
plt.xlabel("Race")
plt.ylabel("Count")
plt.legend(title="Income")
plt.show()
```



```
[37]: # @title Histogram
df.hist(figsize=(12,12), layout=(4,3), sharex=False);
# Noticeable that most of the data of capital gain and loss is 0 so dropping
↳ them
df.drop(['capital-gain', 'capital-loss'], axis=1, inplace=True)
```

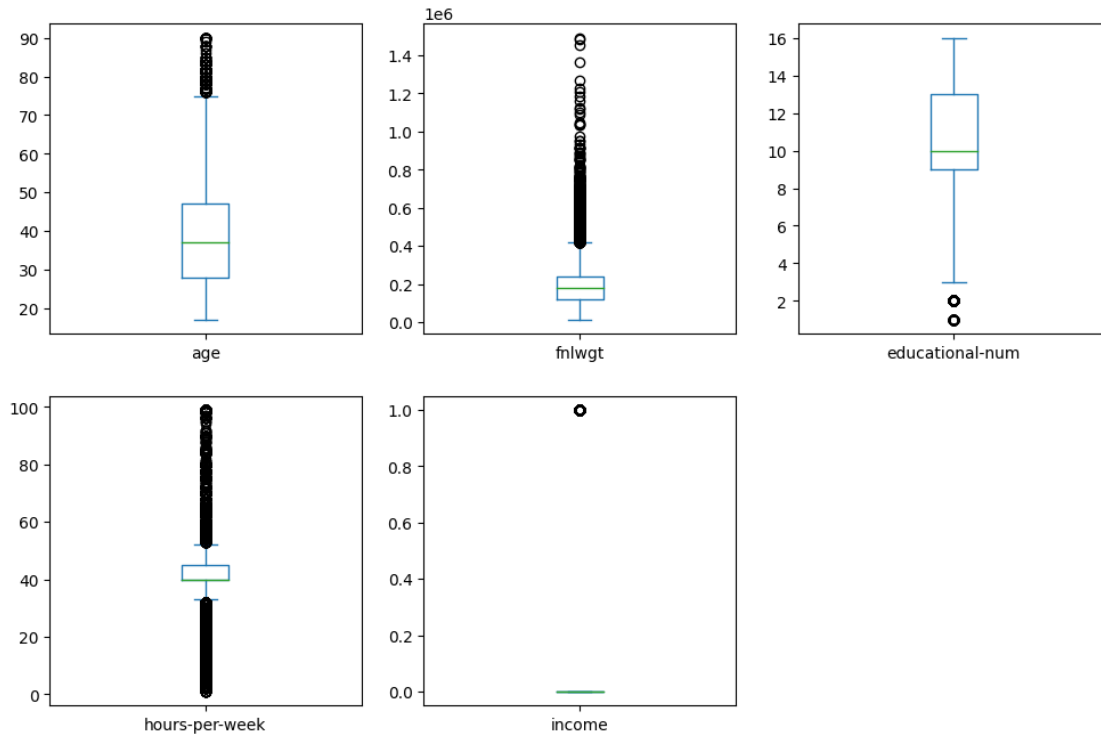


```
[38]: # @title After Dropping capital gain, loss
df.hist(figsize=(12,12), layout=(4,3), sharex=False);
```



```
[39]: # @title Boxplot for checking Outliers

df.plot(kind='box',figsize=(12,12),layout=(3,3),sharex=False,subplots=True);
```



```
[40]: # To find out the Categorical dataset

categorical_df = df.select_dtypes(object)
categorical_df
```

```
[40]:
```

	workclass	education	marital-status	occupation
0	Private	School	Unmarried	Machine-op-inspct
1	Private	High School	Married	Farming-fishing
2	Local-gov	Higher-Education	Married	Protective-serv
3	Private	Higher-Education	Married	Machine-op-inspct
5	Private	School	Unmarried	Other-service
...
48837	Private	Higher-Education	Married	Tech-support
48838	Private	High School	Married	Machine-op-inspct
48839	Private	High School	Single	Adm-clerical
48840	Private	High School	Unmarried	Adm-clerical
48841	Self-emp-inc	High School	Married	Exec-managerial

```

relationship  race  gender  native-country
```

0	Own-child	Black	Male	United-States
1	Husband	White	Male	United-States
2	Husband	White	Male	United-States
3	Husband	Black	Male	United-States
5	Not-in-family	White	Male	United-States
...
48837	Wife	White	Female	United-States
48838	Husband	White	Male	United-States
48839	Unmarried	White	Female	United-States
48840	Own-child	White	Male	United-States
48841	Wife	White	Female	United-States

[45222 rows x 8 columns]

```
[41]: #Checking nan values
df.isna().sum()
```

```
[41]: age          0
workclass        0
fnlwgt          0
education        0
educational-num  0
marital-status   0
occupation       0
relationship     0
race            0
gender          0
hours-per-week   0
native-country   0
income          0
dtype: int64
```

```
[42]: from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
# One-hot encoding categorical variables
df = pd.get_dummies(df, columns=['education', 'marital-status', 'race',
    ↪ 'gender', 'relationship', 'occupation', 'workclass', 'native-country'])

# Splitting the dataset into features and target variable
X = df.drop('income', axis=1)
y = df['income']

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

# Scaling the features
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Prediction Models

```
[43]: # @title K-Nearest Neighbours

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score

# Fitting the KNN model
knn = KNeighborsClassifier(n_neighbors=5) # You can tune n_neighbors
knn.fit(X_train, y_train)

# Predicting the test set results
y_pred_knn = knn.predict(X_test)

# Making the Confusion Matrix
cm_knn = confusion_matrix(y_test, y_pred_knn)
print("KNN Confusion Matrix:")
print(cm_knn)

# Calculating the accuracy score
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print('KNN Accuracy:', accuracy_knn)

# Calculating the accuracy score using the confusion matrix
accuracy_from_cm_knn = (cm_knn[0, 0] + cm_knn[1, 1]) / cm_knn.sum()
print('KNN Accuracy from Confusion Matrix:', accuracy_from_cm_knn)
```

KNN Confusion Matrix:

```
[[6097  745]
 [ 949 1254]]
```

KNN Accuracy: 0.8127142067440575

KNN Accuracy from Confusion Matrix: 0.8127142067440575

```
[44]: # @title Decision Tree

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score

# Fitting the Decision Tree model
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
```

```

# Predicting the test set results
y_pred_dt = dt.predict(X_test)

# Making the Confusion Matrix
cm_dt = confusion_matrix(y_test, y_pred_dt)
print("Decision Tree Confusion Matrix:")
print(cm_dt)

# Calculating the accuracy score
accuracy_dt = accuracy_score(y_test, y_pred_dt)
print('Decision Tree Accuracy:', accuracy_dt)

# Calculating the accuracy score using the confusion matrix
accuracy_from_cm_dt = (cm_dt[0, 0] + cm_dt[1, 1]) / cm_dt.sum()
print('Decision Tree Accuracy from Confusion Matrix:', accuracy_from_cm_dt)

```

Decision Tree Confusion Matrix:

```
[[5786 1056]
 [1018 1185]]
```

Decision Tree Accuracy: 0.770702045328911

Decision Tree Accuracy from Confusion Matrix: 0.770702045328911

```

[45]: # @title Random Forest
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score

# Fitting the Random Forest model
rf = RandomForestClassifier()
rf.fit(X_train, y_train)

# Predicting the test set results
y_pred_rf = rf.predict(X_test)

# Making the Confusion Matrix
cm_rf = confusion_matrix(y_test, y_pred_rf)
print("Random Forest Confusion Matrix:")
print(cm_rf)

# Calculating the accuracy score
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print('Random Forest Accuracy:', accuracy_rf)

# Calculating the accuracy score using the confusion matrix
accuracy_from_cm_rf = (cm_rf[0, 0] + cm_rf[1, 1]) / cm_rf.sum()
print('Random Forest Accuracy from Confusion Matrix:', accuracy_from_cm_rf)

```

Random Forest Confusion Matrix:

```
[[6219  623]
```



```
[ 968 1235]]
```

Random Forest Accuracy: 0.8241017136539525

Random Forest Accuracy from Confusion Matrix: 0.8241017136539525

```
[46]: # @title Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score

# Fitting the Logistic Regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)

# Predicting the test set results
y_pred_logreg = logreg.predict(X_test)

# Making the Confusion Matrix
cm_logreg = confusion_matrix(y_test, y_pred_logreg)
print("Logistic Regression Confusion Matrix:")
print(cm_logreg)

# Calculating the accuracy score using the accuracy_score function
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
print('Logistic Regression Accuracy:', accuracy_logreg)

# Calculating the accuracy score using the confusion matrix
accuracy_from_cm_logreg = (cm_logreg[0, 0] + cm_logreg[1, 1]) / cm_logreg.sum()
print('Logistic Regression Accuracy from Confusion Matrix:',
      accuracy_from_cm_logreg)
```

Logistic Regression Confusion Matrix:

```
[[6305  537]
```

```
 [ 972 1231]]
```

Logistic Regression Accuracy: 0.833167495854063

Logistic Regression Accuracy from Confusion Matrix: 0.833167495854063

```
[47]: # @title Accuracy Check

# Summary of Accuracy Scores

print('K Nearest Neighbour Accuracy:', accuracy_from_cm_knn)
print('Decision Tree Accuracy:', accuracy_from_cm_dt)
print('Random Forest Accuracy:', accuracy_from_cm_rf)
print('Logistic Regression Accuracy:', accuracy_from_cm_logreg)
```

K Nearest Neighbour Accuracy: 0.8127142067440575

Decision Tree Accuracy: 0.770702045328911

Random Forest Accuracy: 0.8241017136539525

Logistic Regression Accuracy: 0.833167495854063